

THE QUESTION EDUCATION CANNOT ANSWER

*Generative AI, the Collapse of Answer-Retrieval, and India's Triple Inequality
of Access, Prompting Literacy and Language*

An Academic Essay for Indian Education

by **Vidya Sharma**

Mentored by Youth Vision India

Prepared for publication by Youth Vision India | October 2026

Acknowledgements

This essay exists because of **Youth Vision India (YVI)**. It began as a private confusion about my own learning and became a researched argument because an organisation run by and for young people took that confusion seriously, and then insisted that I do the work of understanding it properly.

YVI and the cause of education

In a country with roughly 24.69 crore school students and about 1.01 crore teachers, the national conversation about education is often conducted about young people rather than with them. YVI works against that pattern. It treats education as a foundation for wider youth participation and gives young Indians a platform to examine, question and write about the system they are educated in. Its mission is “To empower youth as leaders of change, fostering innovation, integrity, and inclusive growth for a stronger and progressive India.” Its broader mission is to build socially responsible youth leadership through education, community engagement, nation-building initiatives and sustainable development programmes.

That commitment matters particularly now. As this essay argues, artificial intelligence is entering Indian classrooms faster than curricula, teachers and infrastructure can fully absorb it. Youth-led mentorship offers one possible mechanism for developing the critical inquiry that AI-era education increasingly requires.

How YVI promotes education

YVI's contribution to education is practical rather than rhetorical. In my own experience, it shows up in these ways:

- **A platform for youth voices.** By mentoring, editing and now publishing student research, YVI tells young people that their analysis of education policy is worth reading.
- **A standard of evidence.** YVI's mentors consistently pushed me past surface observations toward what the data actually shows, even where it complicated my thesis.
- **Leadership beyond the classroom.** Through its youth leadership and development sessions, YVI gives students opportunities to build communication.
- **Education for health and dignity.** YVI's menstrual-awareness sessions bring conversations around menstrual health and hygiene into spaces where young people can ask questions openly and learn from reliable information. By addressing topics that are often overlooked in conventional classrooms, these sessions broaden what education can mean in practice.

These activities matter to this essay not because they prove a particular position on AI, but because they illustrate a broader educational principle: learning becomes more meaningful when young people are invited to question, engage and apply what they learn.

A Personal Preface

I did not arrive at this topic through a textbook. I arrived at it through a question that embarrassed me. It was late into my preparation for a competitive exam, and I had just spent an hour reading AI-generated notes, confident that I had studied well, only to realise, when I tried to explain the concept to someone else, that I had understood almost nothing. The words had passed through me like water through sand. That experience nagged at me for days. If I, a student who knew I was using AI, could be so easily misled into mistaking fluency for understanding, what would happen to millions of students who did not pause to ask that question?

That is where this essay began. Not with a statistic, but with a moment of honest confusion about my own learning. The research I encountered afterwards, from Wharton, Harvard and Carnegie Mellon, and closer to home from UDISE+ and ASER, suggested that what I had experienced might reflect a broader problem: the ease with which fluent AI output can be mistaken for understanding.

Abstract

Generative AI has entered classrooms faster than many schools can absorb it. This essay argues that its most important educational consequence is not academic dishonesty, but the pressure it places on two assumptions that education has long relied upon: that presenting information can produce learning, and that access to a powerful tool necessarily produces equal outcomes. Drawing on randomized controlled trials, cross-national research, and India-specific evidence from UDISE+ 2024–25 and the Annual Status of Education Report (ASER) 2024, it argues that AI can improve learning when deliberately designed around scaffolding and active engagement, but can also encourage dependence when it substitutes for independent reasoning. In India, these questions intersect with three forms of inequality: infrastructure access, the ability to use AI critically and purposefully, and linguistic inclusion. The essay concludes that AI literacy should be treated as a foundational educational capability, alongside reading and numeracy, while AI systems should be constrained and designed to support rather than replace learning.

India at a glance

Indicator	Figure	Source
Schools / students / teachers (2024–25)	14.71 lakh / 24.69 crore / 1.01 crore	UDISE+ 2024–25
Schools with functional computers	57.2% (2023–24) → 64.7% (2024–25)	UDISE+ 2023–24; 2024–25
Schools with internet	53.9% (2023–24) → 63.5% (2024–25)	UDISE+ 2023–24; 2024–25
Rural 14–16s who can use a smartphone	82.2%	ASER 2024
Of those, used a smartphone for an education-related activity in the reference week	64.4%	ASER 2024
Class 5 children (rural) who can read a Class 2 text	48.8%	ASER 2024

Indicator	Figure	Source
Secondary dropout rate / secondary retention	8.2% / 47.2% (2024–25)	UDISE+ 2024–25
AI & Computational Thinking curriculum	Compulsory Classes III–VIII in CBSE schools from 2026–27	CBSE (2026); MoE (2025)

Table 1. India's education system at a glance, the context for the argument that follows.

I. Introduction: The Question Schools Can No Longer Avoid

Education has always claimed to teach students how to think. Generative AI might be the technology that forces schools to prove it.

Since ChatGPT's release in November 2022, classrooms worldwide have been unsettled, not simply by a new form of cheating, but by an old question that can no longer be deferred: when a machine can produce fluent answers on demand, what exactly is school teaching?

In India the question carries unusual weight. The country has about 24.69 crore school students in 14.71 lakh schools, with about 1.01 crore teachers according to UDISE+ 2024–25. CBSE has also introduced a Computational Thinking and Artificial Intelligence curriculum for Classes III–VIII from the 2026–27 academic session. India is therefore deciding, at enormous scale, not whether AI belongs in education, but how it should enter it.

The central educational challenge is not simply that students can now obtain answers more easily. It is that the skills education ought to cultivate, including questioning, source evaluation, independent reasoning and transfer, are no longer guaranteed to be practised merely because a student completes an assignment. Students with stronger conceptual foundations may be better positioned to use AI as a support for their thinking, while students who cannot evaluate an output may be more vulnerable to treating it as a substitute. The result is not necessarily a dramatic new divide. It can be a quieter one, in which the same tool produces different educational outcomes for different students.

Schools that treat Generative AI only as a cheating problem therefore risk missing a deeper structural question: what is education for when information retrieval is nearly frictionless?

II. The Collapse of Answer-Retrieval as a Learning Objective

For decades, Bloom's Taxonomy has provided a useful framework for distinguishing foundational forms of learning from higher-order capacities such as analysis, evaluation and creation. Generative AI complicates the relationship between these categories because a model can now produce outputs that resemble higher-order work without the student necessarily developing the underlying ability.

A peer-reviewed study in the *Journal of Marketing Education* examined Generative AI through the lens of Bloom's Taxonomy and argued that AI assistance may be more straightforwardly useful for lower-order cognitive tasks than for the higher-order educational work of evaluation and creation (Gonsalves, 2024). The important distinction is not that AI is incapable of analysis or creation. It is that producing an answer is not equivalent to a student learning how to perform the reasoning that produced it.

This distinction matters acutely in India, where high-stakes examinations such as board examinations, JEE, NEET and CUET place substantial value on what students can independently recall, apply and solve under time pressure. If assessment rewards outputs that AI can increasingly produce, the credibility of those assessments will depend more heavily on whether students can demonstrate judgement and process as well as answers.

Figure 1 — Bloom’s Taxonomy: Where AI Assists vs. Where Education Must Lead*Relative AI capability across cognitive levels (illustrative, based on Gonsalves 2024 and Zhai et al. 2025)*

Remembering		Very High
Understanding		High
Applying		Moderate
Analysing		Mixed
Evaluating		Limited
Creating		Minimal

Figure 1. AI capability declines sharply at higher cognitive levels, precisely where genuine intellectual development lies. Navy bars show tasks AI performs well; saffron bars show where human learning remains irreplaceable.

This is not surprising from the perspective of cognitive science. Robert A. Bjork and Elizabeth L. Bjork’s research on ‘desirable difficulties’ demonstrated that conditions which slow short-term performance, such as spaced repetition, interleaved practice and delayed retrieval, produce substantially stronger long-term retention and transfer (Bjork & Bjork, 1994, pp. 475–476). Learning that feels easy in the moment is frequently not learning that lasts.

Generative AI introduces a profound paradox here. It produces answers with almost no friction, and that effortlessness generates what researchers call the ‘illusion of fluency’: the feeling of having understood something one has merely read. A student who submits an AI-generated answer may have scored highly but has bypassed the cognitive effort that makes understanding durable. For a Class 12 student summarising a chapter for the boards, this looks like efficiency. In many cases it is a loss of depth that will only become apparent when the scaffolding is removed.

The most rigorous experimental evidence comes from a large randomised controlled trial by researchers at the Wharton School, University of Pennsylvania, published in the *Proceedings of the National Academy of Sciences*. Nearly one thousand high school mathematics students were given access to AI tutors. Students using unrestricted GPT-4 **improved performance by 48%** during AI-assisted practice, but **performed 17% worse** on unassisted examinations than the control group, a reversal that points directly to reduced skill acquisition (Bastani et al., 2025, *PNAS*, 122(26), e2422633122).

LANDMARK RCT FINDING

Students using unrestricted GPT-4 improved performance during practice by 48%, but scored 17% worse on unassisted exams compared to the control group, a reversal consistent with reduced skill acquisition. — Bastani et al. (2025), *Proceedings of the National Academy of Sciences*

Critically, the same study showed that design matters enormously. When the AI tutor was constrained to provide hints rather than full answers, the negative learning effects were substantially mitigated. The problem is not AI itself; it is how AI is deployed.

A contrasting but complementary study from Harvard University’s Department of Physics, published in *Scientific Reports* (Kestin, Miller, Klales, Milbourne & Ponti, 2025), found that a carefully designed AI tutor produced learning gains more than twice as large as those of a

human-led active-learning classroom, and in less time. Students also reported greater motivation. The operative phrase is ‘carefully designed’. The Harvard team built their tutor around established pedagogical principles: brevity, guided scaffolding, hallucination prevention through curated answer keys, and Socratic dialogue rather than answer delivery. The difference between the Wharton and Harvard findings is not the technology. It is the pedagogy.




Figure 2. Unrestricted AI boosts in-session performance but harms long-term learning. Carefully designed AI tutoring can double learning gains. The difference is design, not technology.

The Stanford Institute for Human-Centered Artificial Intelligence’s 2025 AI Index Report crystallises the institutional challenge: while 81% of US K–12 computer science teachers believe AI should be part of foundational education, fewer than half feel prepared to teach it (Stanford HAI, 2025). According to a nationally representative 2025 RAND survey, 54% of middle and high school students already use AI for school, but more than 80% reported that their teachers had never explicitly taught them how to use it (Stanford SCALE, 2026). The tools have arrived. The pedagogy has not. India’s own readiness figures, examined in Section IV, are considerably starker.

III. The Prompting Literacy Gap: A New Inequality

It is tempting to view Generative AI as a democratising force, a free, infinitely patient tutor available to any student with a smartphone. That framing is appealing and not entirely wrong. But it obscures a mechanism through which AI is quietly amplifying existing inequalities rather than dissolving them.

A systematic review of 51 studies on ChatGPT in education, published in *Frontiers in Education*, identified a consistent pattern: students with stronger prior knowledge and more developed critical thinking skills derive substantially greater benefit from AI, while novices frequently fail to detect errors or misleading outputs (Frontiers in Education, 2024). If the primary beneficiaries of AI tools are students who are already ahead, the tool does not level the playing field. It tilts it further.

This pattern extends beyond school. A 2025 working paper from researchers at Harvard Business School and Stanford University, *The GenAI Wall Effect*, studied 78 workers given identical access to the same AI tools. While AI improved performance on abstract, generative tasks such as brainstorming and drafting, it provided almost no benefit to novices on tasks requiring contextual judgement. The researchers’ conclusion applies directly to education: ‘GenAI can provide the map, but navigating the terrain still requires experience’ (Bojinov et al., 2025).

A 2025 study from Microsoft Research and Carnegie Mellon University reinforces this. Analysing 936 real-world AI use cases, researchers found that higher confidence in AI was associated with less critical thinking, while confidence in one's own knowledge was associated with more critical engagement (Lee et al., 2025). When students trust AI without understanding its outputs, they stop questioning, and that cessation of questioning is precisely where learning ends.

KEY FINDING — THE CRITICAL THINKING PARADOX

Higher confidence in AI was linked to less critical thinking. Confidence in one's own abilities led to more. When people trust AI without understanding it, they stop questioning it altogether. — Lee et al. (2025), Microsoft Research and Carnegie Mellon University

Research from the British Educational Research Association adds an institutional dimension: wealthier schools tend to use AI as a supplement to strong existing teaching, while under-resourced schools are more likely to use it as a substitute for what they lack (Allison, 2024, BERA). Evidence from the Brookings Institution, examining remote learning across Ghana, Indonesia, Poland, Saudi Arabia and the United States, found that even where access to technology was widespread, learning gaps between richer and poorer students persisted (Brookings, 2023). Access is a necessary but wholly insufficient condition for educational equity. As the next section shows, India's own data tell the same story.

IV. India's Triple Inequality: Infrastructure, Literacy and Language

The prompting literacy gap is a global phenomenon. But India presents a uniquely compounded version of it, one that combines three distinct inequalities in a way no broad cross-national comparison fully captures.

4.1 The Infrastructure Gap

Before asking whether Indian students can use AI well, one must ask whether they can access it at all in school. Progress is real but incomplete. UDISE+ 2024–25 records that 64.7% of schools now have functional computers and 63.5% have internet access, up from 57.2% and 53.9% a year earlier (UDISE+ 2023–24; Mehta, 2025). Yet this still means that more than one in three schools has no functional computer or internet connection, and for these schools an AI curriculum is an abstraction.

Teacher distribution compounds the problem. UDISE+ 2024–25 counts 1,04,125 single-teacher schools, about 7% of all schools, enrolling roughly 3.38 million students (Mehta, 2025). A single teacher responsible for several grades cannot be expected to supervise, verify and contextualise AI use for each of them. The Azim Premji Foundation's 2021 study found that a majority of children in the surveyed states had limited or no access to online learning during the pandemic (Azim Premji Foundation, 2021), a reminder of how quickly digital dependence becomes educational exclusion.

Figure 3 — Digital Infrastructure in Indian Schools

% of schools with functional computers and internet access; UDISE+ 2023–24 and 2024–25

Schools with computers (2023–24)		57.2%
Schools with computers (2024–25)		64.7%
Schools with internet (2023–24)		53.9%
Schools with internet (2024–25)		63.5%

Figure 3. Digital infrastructure is improving but over a third of Indian schools remain without computers or internet. Sources: UDISE+ 2023–24; Mehta (2025) analysis of UDISE+ 2024–25.

4.2 The Prompting Literacy Gap in the Indian Context

The most revealing Indian evidence on prompting literacy comes not from schools but from households. ASER 2024, which surveyed 649,491 children in 17,997 villages across 605 rural districts, found that access to smartphones among 14–16-year-olds is close to universal: about 90% of both girls and boys report a smartphone at home, and 82.2% know how to use one. But only 57% had used a smartphone for an educational activity in the preceding week, while 76% had used it for social media (ASER Centre, 2025; PIB, 2025).


Figure 4 — Smartphones in Rural India: Access vs. Educational Use (ages 14–16) % of children aged 14–16; ASER 2024 (rural India)		
Have a smartphone at home		~90%
Know how to use a smartphone		82.2%
Used for social media (past week)		76%
Used for education (past week)		57%
Boys owning a smartphone		36.2%
Girls owning a smartphone		26.9%

Figure 4. Access is no longer the main barrier for rural teenagers; the gap is in purposeful, educational use, and girls are less likely to own a device. Source: ASER Centre (2025).

This is prompting the literacy gap in its Indian form. A device in the hand does not produce a critical user. The foundations beneath are also fragile: in ASER 2024, only 48.8% of rural Class 5 children could read a Class 2 text, and in government schools just 23.4% of Class 3 children could (ASER Centre, 2025). A child who cannot fluently read a short passage cannot reliably detect whether a confident AI paragraph is correct. The most-cited finding in this essay, that novices are the ones who fail to detect AI errors, therefore describes not a niche group but a large share of India's school population.

The retention data deepen the concern. UDISE+ 2024–25 records a secondary dropout rate of 8.2%, down from 10.9% the previous year, but retention at the secondary level stands at only 47.2% (UDISE+ 2024–25, as reported by Careers360, 2025). ASER 2024 found that 7.9% of 15–16-year-olds in rural India are not enrolled in any school, down sharply from 13.1% in 2018. These are the students for whom AI's democratising potential is most often invoked, yet who face the steepest structural barriers to realising it.

Total enrolment from pre-primary to Class 12 has fallen from 26.52 crore in 2021–22 to 24.69 crore in 2024–25, a decline of about 1.83 crore, or roughly 6.9% (UDISE+, 2021–22 to 2024–25; Mehta, 2025). Analysts attribute this to demographic transition, a post-pandemic shift between school types and improved data quality through Aadhaar-linked student records rather than to a single cause. AI will not solve a challenge of this magnitude. It can deepen it, or

it can be one tool among many used thoughtfully to address it. Which outcome materialises depends on policy choices currently being made.

4.3 The Language Gap: A Dimension Unique to India

There is a third inequality that the global literature rarely addresses with adequate seriousness, and which India faces with unique intensity: the linguistic exclusion embedded in most major AI systems.

UNESCO's 2025 State of Education Report for India, *Bhasha Matters*, documents that India is home to 1,369 mother tongues and emphasises that instruction in a child's first language is foundational to comprehension, cognitive development and academic success (UNESCO, 2025). NEP 2020 mandates instruction in the learner's home language through at least Grade 5 and preferably through Grade 8. Yet most leading AI systems were trained predominantly on English-language data and perform measurably worse in Indian languages. Recent benchmark research found that frontier multimodal AI models regressed by an average of **8.3 points** from English to Hindi and **5.5 points** from English to Telugu on equivalent visual question-answering tasks (HinTel-AlignBench, 2024). For students whose primary learning language is Odia, Bhojpuri, Maithili or any of dozens of other mother tongues, the gap is almost certainly larger.

The Government of India has recognised this through the Bhashini initiative, an AI-powered translation and speech platform targeting all 22 scheduled languages, and AICTE's Anuvadini app, which translates educational texts into Indian languages (Digital India, 2025). These efforts matter, but they also show how much work remains before AI tools can be called genuinely inclusive.

Dimension	What the gap looks like	Key statistic	Source
Infrastructure	No functional computer or internet at school	35.3% of schools lack computers; 36.5% lack internet	UDISE+ 2024–25
Prompting literacy	Access without purposeful, critical use	82.2% can use a smartphone; 57% used it for education	ASER 2024
Foundational skills	Too weak to verify AI output	48.8% of rural Class 5 read a Class 2 text	ASER 2024
Language inclusion	AI underperforms in most Indian languages	–8.3 pts English→Hindi on benchmarks	HinTel-AlignBench, 2024

Table 2. India faces a compounding inequality across interconnected dimensions. Addressing one without the others risks reproducing the existing hierarchy in digital form. (Lack-of-computer and lack-of-internet shares are the complements of the UDISE+ 2024–25 figures.)

4.4 Policy Ambition and the Readiness Gap

India's policy response is moving fast. NITI Aayog's 2018 National Strategy for Artificial Intelligence named education as a priority sector (NITI Aayog, 2018). On 30 October 2025 the Ministry of Education announced that Artificial Intelligence and Computational Thinking would be introduced from Class 3 from the 2026–27 session under NEP 2020 and the National Curriculum Framework for School Education 2023, with teacher training delivered through NISHTHA. CBSE has since issued a curriculum framework making the subject part of Classes III to VIII in the 2026–27 session (CBSE, 2026). The School Education Secretary has acknowledged the central challenge: reaching more than one crore teachers (Careers360, 2025).

Before this rollout, over 18,000 CBSE schools already offered AI as a skill subject from Class 6 through a 15-hour module, but CBSE schools are a minority of India's 14.71 lakh schools.

Two states offer instructive early case studies. Punjab became the first Indian state to roll out an AI curriculum across government schools, and Odisha's AI Policy 2025 sets a target of AI instruction in 35% of schools by 2029 and 90% by 2036. Yet experts estimate that training five million teachers in a first phase alone would require approximately 200–500 core AI experts, 10,000 master trainers and 600,000 district-level trainers, a capacity India does not currently have (The Secretariat, 2026).

A policy decision made in Delhi does not remain a policy decision. It becomes a daily classroom reality for a city school in Mumbai and a single-teacher primary school in rural Bihar, often on very different terms. The government school sector alone enrolls 12.16 crore students, 49.2% of the total (Mehta, 2025), and it is here that the gap between ambition and readiness is widest.

V. The Counterargument: AI as Democratic Access to Knowledge

The strongest challenge to this essay's argument is the democratisation thesis: the claim that Generative AI gives every student, regardless of geography or economic background, access to a knowledgeable, patient, always-available tutor. In India, where classrooms are overcrowded, teachers are scarce and private tuition is unaffordable for many families, this is not a minor claim. ASER 2024 itself records that around 33% of rural children aged 6–14 attend private schools and that private tutoring remains widespread; for families priced out of both, a free AI tutor could matter enormously. It deserves serious engagement.

There is credible experimental evidence in its support. The Kestin et al. (2025) study at Harvard produced learning gains more than twice as large as those of skilled in-person instruction. A pre-registered experiment by Lira et al. found that participants who practised writing with structured AI assistance outperformed those who practised without it on subsequent independent tasks, suggesting genuine skill transfer (Lira et al., 2024). A 2025 meta-analysis in *npj Science of Learning*, examining 28 studies involving nearly 4,600 students, found that AI tutoring systems outperformed traditional teaching when they combined personalised pacing, immediate feedback and step-by-step guidance (npj Science of Learning, 2025).

The key phrase in all of this evidence is 'carefully designed'. The Harvard study used a custom tutor with hallucination prevention and explicit Socratic constraints, not a general-purpose chatbot. The Lira et al. study observed benefits when AI provided models for learning, not when it simply generated answers. The meta-analysis found benefits under conditions of adaptive feedback and scaffolding.

The democratisation thesis holds, but only under structured conditions, with well-designed tools, and for students who already have enough foundational knowledge to evaluate what they receive. The third condition is the most limiting. Research in *npj Digital Medicine* on novice medical students found that incorrect AI explanations significantly reduced diagnostic accuracy, while correct explanations provided no meaningful advantage over having no AI explanation at all (npj Digital Medicine study, cited in Bastani et al., 2025).

For students who cannot reliably identify when an AI is wrong, which in India includes the many whose foundational reading and numeracy are still developing, access to Generative AI does not democratise knowledge. It democratises exposure to confident-sounding error.

VI. What the Evidence Demands: A Framework for Action

If the evidence above is taken seriously, the question is not whether to include AI in education. That decision has effectively been made by millions of students, their smartphones and the market, and now by the curriculum itself. The question is what Indian education must do to ensure that AI deepens learning rather than replacing it, and narrows inequality rather than entrenching it.

Priority	What it means for India	Evidence basis
Treat AI literacy as a foundational skill	Teach students to question, verify and critically evaluate AI outputs, as core curriculum alongside reading and numeracy	Lee et al. (2025); Frontiers in Education (2024); CBSE (2026)
Constrain, don't ban	Design AI interactions around hints, scaffolding and guided reasoning, not answer delivery	Bastani et al. (2025)
Redesign assessment	Shift board and entrance evaluation toward judgement, questioning and process, not recall	Bjork & Bjork (1994); Gonsalves (2024)
Close the infrastructure gap first	Prioritise the 35% of schools without computers and the 36% without internet, and single-teacher schools	UDISE+ 2024–25
Fix foundations alongside AI	Sustain NIPUN Bharat-style foundational literacy so students can evaluate what AI tells them	ASER 2024
Train teachers before students	Teacher preparedness is the largest bottleneck: over one crore teachers to orient	Stanford HAI (2025); Careers360 (2025)
Build for linguistic inclusion	AI tools for Indian students must work in Indian languages	UNESCO (2025); Bhashini
Mobilise youth-led networks	Use youth organisations for peer mentoring and AI-literacy workshops where formal capacity is thin	This essay's argument; see Acknowledgements

Table 3. A synthesis of evidence-based priorities. None is optional; each addresses a distinct failure mode the research has identified.

The final priority deserves a word. Formal systems move slowly, and the Class 3 rollout is happening faster than teacher training can keep pace. Youth-led organisations can close part of that gap. They can run peer-led workshops on verifying AI outputs, mentor students through evidence-based writing as YVI does, and reach the students and communities that large programmes reach last. They are not a substitute for investment in schools, but they are a multiplier for it.

VII. Conclusion: Redesigning the Question, Not Just the Tool


The most important impact of Generative AI on education is not that students are using it to cheat. It is that AI has made the purpose of education impossible to avoid confronting. When a

machine can retrieve answers, memorise content and produce polished written arguments on demand, it exposes, with uncomfortable clarity, what education was measuring all along. What remains irreplaceable are the capacities AI cannot perform on a student's behalf: the ability to form better questions, evaluate competing claims with rigour, direct inquiry with intellectual honesty, and build the foundational knowledge that makes all of this possible.

For India, the stakes are uniquely high. A system of 24.69 crore students is simultaneously rolling out a compulsory AI curriculum from Class 3 and managing a situation in which over a third of schools lack a functional computer or internet connection, fewer than half of rural Class 5 children can read a Class 2 text, and most AI systems handle their mother tongues poorly. India's teenagers already hold the device; what they lack is the literacy to use it well. The gap between ambition and readiness is not a minor implementation challenge. It is the central policy question of the next decade.

The institutions that will genuinely serve their students in an AI-saturated world are those that treat AI literacy as foundational, alongside reading and numeracy, not as a supplement for those already ahead. They redesign assessment around judgement and process rather than information recall. And they recognise that education is no longer about producing answers, but about forming people who know which answers are worth trusting and why.

Generative AI did not create this challenge. It simply made it impossible to ignore any longer.



Bibliography

Journal Articles and Books

- Allison, S. (2024). AI in schools: Equity implications of uneven adoption. British Educational Research Association (BERA). <https://www.bera.ac.uk>
- Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakçı, Ö., & Mariman, R. (2025). Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*, 122(26), e2422633122. <https://doi.org/10.1073/pnas.2422633122>
- Bjork, R. A., & Bjork, E. L. (1994). A new theory of disuse and an old theory of stimulus fluctuation. In A. Healy, S. Kosslyn, & R. Shiffrin (Eds.), *From learning processes to cognitive processes: Essays in honor of William K. Estes* (Vol. 2, pp. 475–476). Erlbaum.
- Bloom, B. S., Engelhart, M. D., Furst, E. J., Hill, W. H., & Krathwohl, D. R. (1956). *Taxonomy of educational objectives, Handbook I: The cognitive domain*. David McKay.
- Bojinov, I., Dell'Acqua, F., Teodoridis, F., & Edelman, B. (2025). The GenAI wall effect. Harvard Business School & Stanford University Working Paper.
- Frontiers in Education. (2024). A systematic review of ChatGPT in education across 51 studies. <https://www.frontiersin.org>
- Gonsalves, A. (2024). Generative AI and Bloom's Taxonomy: Implications for marketing education. *Journal of Marketing Education*. <https://doi.org/10.1177/02734753241236456>
- Kestin, G., Miller, K., Klales, A., Milbourne, T., & Ponti, G. (2025). AI tutoring outperforms in-class active learning: An RCT introducing a novel research-based design in an authentic educational setting. *Scientific Reports*, 15, Article 17458. <https://doi.org/10.1038/s41598-025-97652-6>
- Lee, M., Sellen, A., Rothschild, D. M., & colleagues. (2025). Critical thinking and AI confidence: Analyzing 936 real-world cases. Microsoft Research & Carnegie Mellon University.
- Lira, C., & colleagues. (2024). Writing with AI assistance: A pre-registered experiment on learning transfer. University of Pennsylvania & Harvard University.

Reports, Policy Documents and Data Sources (India)

- ASER Centre. (2025). *Annual Status of Education Report (Rural) 2024*. Pratham Foundation. <https://asercentre.org>
- Azim Premji Foundation. (2021). Loss of learning during the pandemic. <https://azimpremjiuniversity.edu.in/apf-reports>
- Central Board of Secondary Education. (2026). *CBSE curriculum framework for Computational Thinking and Artificial Intelligence, Classes III–VIII (Session 2026–27)* (Circular Acad-15/2026, 1 April 2026). <https://cbseacademic.nic.in>
- Department of School Education and Literacy, Ministry of Education, Government of India. (2024). *UDISE+ 2023–24 Flash Statistics*. <https://udiseplus.gov.in>
- Department of School Education and Literacy, Ministry of Education, Government of India. (2025). *UDISE+ 2024–25 report*. <https://udiseplus.gov.in>
- Ministry of Education, Government of India. (2025, October 30). *AI and Computational Thinking to be introduced from Class 3 from 2026–27* [Press release]. Press Information Bureau. <https://pib.gov.in>
- Ministry of Electronics and Information Technology, Government of India. (2025). Bhashini: Digital India language initiative. <https://bhashini.gov.in>
- NITI Aayog. (2018). *National Strategy for Artificial Intelligence*. <https://niti.gov.in>
- Press Information Bureau. (2025). *Annual Status of Education Report (ASER) 2024* [Fact sheet]. <https://pib.gov.in>
- UNESCO. (2025). *Bhasha Matters: State of the Education Report for India 2025 — Mother tongue and multilingual education*. UNESCO Regional Office for South Asia. <https://unesdoc.unesco.org>

Reports, Policy Documents and Data Sources (International)

- Brookings Institution. (2023). Remote learning, equity, and the digital divide: Evidence from five countries. <https://www.brookings.edu>
- ICRIER. (2025, January). AI in school education: Towards a preparedness framework. ICRIER Policy Brief. https://icrier.org/pdf/AI_in_School_Education.pdf
- Stanford HAI. (2025). *AI Index Report 2025*. Stanford Institute for Human-Centered Artificial Intelligence. <https://hai.stanford.edu/ai-index/2025-ai-index-report>
- Stanford SCALE. (2026). Artificial intelligence in K–12 schools. Stanford Center for Learning & Equity. <https://scale.stanford.edu>
- Zhai, X., Wibowo, S., & Li, L. D. (2024). The effects of over-reliance on AI dialogue systems on students' cognitive abilities: A systematic review. *Smart Learning Environments*, 11(1), 28. <https://doi.org/10.1186/s40561-024-00316-7>

News and Secondary Sources

- Careers360. (2025). *Centre to introduce AI in school curriculum from Class 3 onwards from 2026–27 session; UDISE+ 2024: School enrolment drops by 11 lakh with secondary retention at 47%; teachers cross 1 crore*. <https://news.careers360.com>
- Deccan Herald / The Diplomat. (2025, November). The risk of unequal AI education in India's schools. <https://www.deccanherald.com>
- Mehta, A. C. (2025). *Status of school education in India: An analysis based on UDISE+ 2021–22 to 2024–25*. Education for All in India. <https://educationforallinindia.com>
- NORRAG. (2024, October). Digitalisation in education: Can AI bridge India's digital divide? <https://norrageducation.org>
- The Secretariat. (2026, March). India's schools step into the AI era with nationwide training plan. <https://thesecretariat.in>

